

Sr. No.	Practical	Signature
1	To implement various descriptive statistics methods 1.1) central tendency, quartile and interquartile 1.2) univariate, bivariate and multivariate.	
2	To implement data cleaning 2.1) Removing leading or lagging spaces from a data entry 2.2) Removing nonprintable characters from a data entry 2.3) Data cleaning: handling missing values, type conversion, data transformations, removing duplicates. 2.4) To detect outliers in the given data.	
3	Regression Analysis 3.1) To perform regression analysis using single linear regression. 3.2) To perform regression analysis using multiple linear regression. 3.3) To perform logistic regression analysis	
4	Classification 4.1) To implement classification using decision tree induction 4.2) To implement classification using Naïve Bayes algorithm 4.3) To implement classification using decision tree induction with various attribute selection methods(Information Gain, Gini index and Gain ratio)	
5	Clustering Algorithm 5.1) To implement clustering using K-Means Algorithm 5.2) To perform hierarchical clustering	
6	To implement PCA (Principal Component Analysis).	
7	To explore the given data and identify the patterns in it.	
8	8.1) To evaluate binary classification model using confusion matrix along with precision and recall. 8.2) To evaluate multi-class classification model using confusion matrix along with precision and recall.	

9.	Use an appropriate dataset and create a supervised learning model, Analyse the model with ROC-AUC.	
10.	Consider a case study problem and implement an appropriate model and evaluate it.	
11.	Write a program to implement 11.1 Bagging and boosting model. 11.2 Cross validation methods	

Aim: To implement various descriptive statistics methods

1.1)central tendency, quartile and interquartile

```
import pandas as pd

import numpy as np

# Ensure all columns have exactly 10 entries
Data = {
    'Student_id': [101, 102, 103, 104, 105, 106, 207, 108, 109, 110],
    'Age': [18, 19, 18, 20, 19, 21, 18, 20, 18, 22],
    'score': [85, 59, 27, 89, 58, 87, 58, 90, 82, 89],
    'Study_hours': [5, 7, 4, 8, 6, 3, 7, 5, 6, 9]
}
df = pd.DataFrame(Data)
print("Original DataFrame:")
print(df)
print("\n")
print("Descriptive Statistics using .describe():")
print(df.describe())
print("\n")
print("Individual Statistical Measures:")

# Central tendency
print(f"Mean of score: {df['score'].mean():.2f}")
print(f"Median of score: {df['score'].median():.2f}")
print(f"Mode of Age: {df['Age'].mode().tolist()}")

# Quartiles
print(f"25th percentile (Q1) of score: {df['score'].quantile(0.25):.2f}")
print(f"50th percentile (Q2 / median) of score: {df['score'].quantile(0.50):.2f}")
print(f"75th percentile (Q3) of score: {df['score'].quantile(0.75):.2f}")

# Interquartile Range
iqr_score = df['score'].quantile(0.75) - df['score'].quantile(0.25)
print(f"Interquartile Range (IQR) of score: {iqr_score:.2f}")
```

OUTPUT:

```
Original DataFrame:
  Student_id  Age  score  Study_hours
0         101   18    85           5
1         102   19    59           7
2         103   18    27           4
3         104   20    89           8
4         105   19    58           6
5         106   21    87           3
6         207   18    58           7
7         108   20    90           5
8         109   18    82           6
9         110   22    89           9

Descriptive Statistics using .describe():
count      10.000000  10.000000  10.000000  10.000000
mean       115.500000  19.300000  72.400000   6.000000
std        32.287769   1.418136  21.030137   1.825742
min        101.000000  18.000000  27.000000   3.000000
25%        103.250000  18.000000  58.250000   5.000000
50%        105.500000  19.000000  83.500000   6.000000
75%        108.750000  20.000000  88.500000   7.000000
max        207.000000  22.000000  90.000000   9.000000

Individual Statistical Measures:
Mean of score: 72.40
Median of score: 83.50
Mode of Age: [18]
25th percentile (Q1) of score: 58.25
50th percentile (Q2 / median) of score: 83.50
75th percentile (Q3) of score: 88.50
Interquartile Range (IQR) of score: 30.25
```

Aim: To implement various descriptive statistics methods

1.2) univariate, bivariate and multivariate.

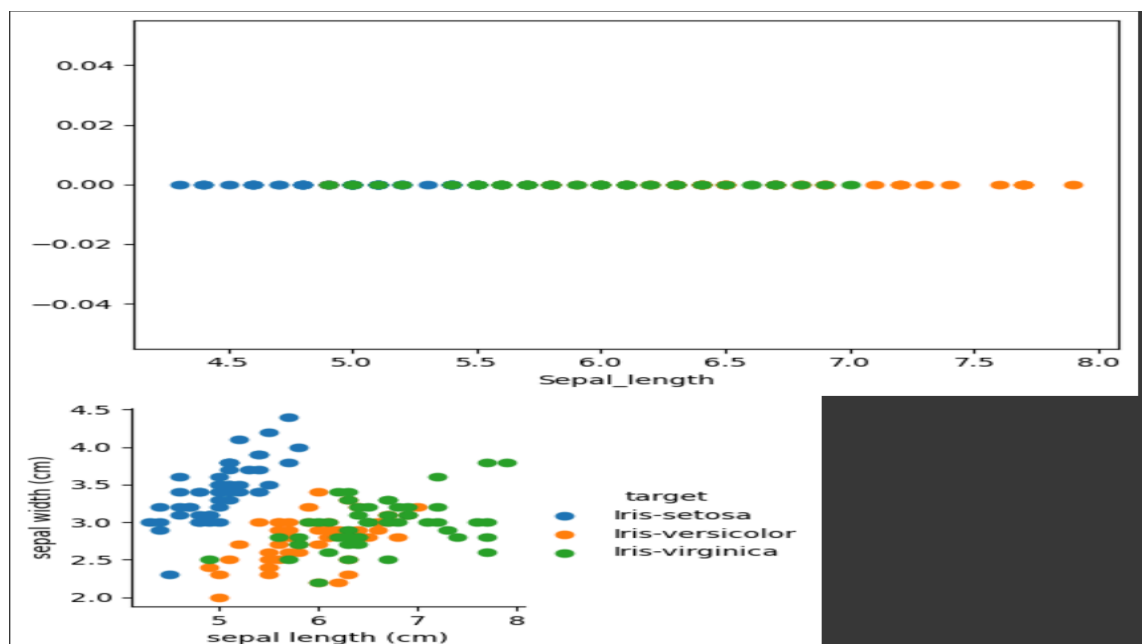
```
import pandas as pd
import numpy as np
import matplotlib.pyplot as plt
import seaborn as sns
df=pd.read_csv("/content/Iris.csv")

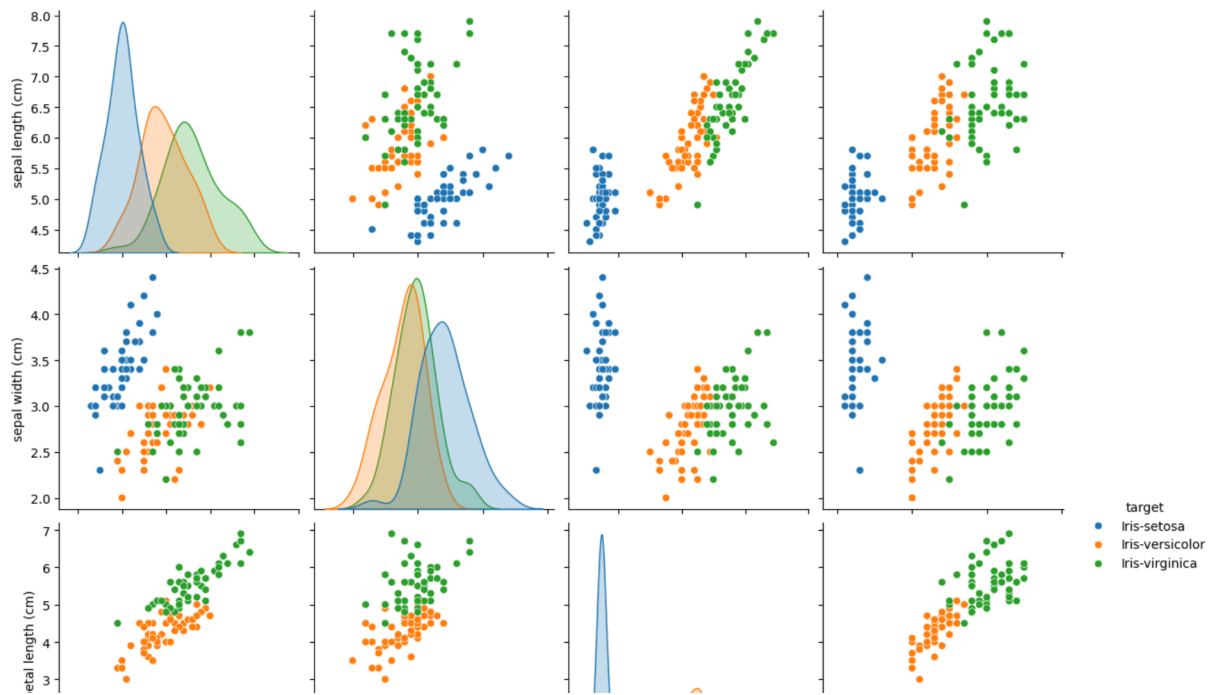
#univariate
df_setosa=df.loc[df['Species']=='Iris-setosa']
df_virginica=df.loc[df['Species']=='Iris-virginica']
df_versicolor=df.loc[df['Species']=='Iris-versicolor']
plt.plot(df_setosa['Sepal_length'],np.zeros_like(df_setosa['Sepal_length']),'o')
plt.plot(df_virginica['Sepal_length'],np.zeros_like(df_virginica['Sepal_length']),'o')
plt.plot(df_versicolor['Sepal_length'],np.zeros_like(df_versicolor['Sepal_length']),'o')
plt.xlabel("Sepal_length")
plt.show()

#bivariate
sns.FacetGrid(df,hue='Species').map(plt.scatter,'Sepal_length','Sepal_width').add_legend()
plt.show()

#Multivariate
sns.pairplot(df,hue='Species',size=3)
plt.show()
```

OUTPUT:





Aim: To implement data cleaning

2.1) Removing leading or lagging spaces from a data entry

```
# Create a sample DataFrame with leading/trailing spaces
data = {'TextColumn': [' hello ', 'world ', ' example']}
df = pd.DataFrame(data)
print("Original DataFrame:")
print(df)

# Remove trailing spaces from the 'TextColumn'
df['TextColumn'] = df['TextColumn'].str.rstrip()

print("\nDataFrame after removing trailing spaces:")
print(df)

# Remove leading spaces from the 'TextColumn'
df['TextColumn'] = df['TextColumn'].str.lstrip()

print("\nDataFrame after removing leading spaces:")
print(df)
```

OUTPUT:

```
Original DataFrame:
  TextColumn
0    hello
1    world
2  example

DataFrame after removing trailing spaces:
  TextColumn
0    hello
1    world
2  example

DataFrame after removing leading spaces:
  TextColumn
0    hello
1    world
2  example
```

2.2) Removing nonprintable characters from a data entry

```
import string
# Create a set of printable characters
printable = set(string.printable)

# Sample data with nonprintable characters
data_with_nonprintable = "This is a string with \n a newline and \r a carriage return."

print("Original string:")
print(data_with_nonprintable)

# Remove nonprintable characters using a list comprehension
cleaned_data = "".join([char for char in data_with_nonprintable if char in printable])

print("\nString after removing nonprintable characters (simpler version):")
print(cleaned_data)
```

OUTPUT:

```
Original string:
This is a string with
  a carriage return.

String after removing nonprintable characters (simpler version):
This is a string with
  a carriage return.
```


2.3) Data cleaning: handling missing values, type conversion, data transformations, removing duplicates.

```
import pandas as pd
import numpy as np
def clean_dataset(df):
    print("-----Handling missing value-----")
    print("Missing values before cleaning:\n", df.isnull().sum())

    # 1. Fill numeric missing values
    for col in df.select_dtypes(include=np.number).columns:
        if df[col].isnull().any():
            df[col] = df[col].fillna(df[col].mean())

    # Fill categorical missing values
    for col in df.select_dtypes(include='object').columns:
        if df[col].isnull().any():
            df[col] = df[col].fillna(df[col].mode()[0])

    print("Missing values after filling:\n", df.isnull().sum())

    # 2. Type conversion
    print("-----Type conversion-----")
    if 'sone_numeric_column_string' in df.columns:
        df['sone_numeric_column_string'] = pd.to_numeric(
            df['sone_numeric_column_string'], errors='coerce')
        df['sone_numeric_column_string'] = df['sone_numeric_column_string'].fillna(
            df['sone_numeric_column_string'].mean())
        print("Converted 'sone_numeric_column_string' to numeric")

    if 'date_column' in df.columns:
        df['date_column'] = pd.to_datetime(df['date_column'], errors='coerce')
        print("Converted 'date_column' to datetime")

    # 3. Data transformation
    print("-----Data transformation-----")
    if 'column_a' in df.columns and 'column_b' in df.columns:
        df['new_feature'] = df['column_a'] * df['column_b']
        print("Created new feature by multiplying 'column_a' and 'column_b'")

    # 4. Removing duplicates
    print("-----Removing duplicates-----")
    initial_rows = len(df)
    df.drop_duplicates(inplace=True)
    print(f"Removed {initial_rows - len(df)} duplicate rows")

    return df
```

```

if __name__ == '__main__':
    data = {
        'numerical_col_1': [1, 2, np.nan, 4, 5],
        'numerical_col_2': [10.5, 11.5, 10.8, np.nan, 12.1],
        'categorical_col': ['A', 'B', 'A', 'C', np.nan],
        'sone_numeric_column_string': ['100', '200', 'abc', '400', '500'],
        'date_column': ['2023-01-01', '2023-01-02', 'invalid date', '2023-01-04', '2023-01-05'],
        'column_a': [1, 2, 3, 4, 5],
        'column_b': [5, 4, 3, 2, 1],
    }

    sample_df = pd.DataFrame(data)
    print("Original dataframe:\n", sample_df)

    cleaned_df = clean_dataset(sample_df.copy())
    print("\nCleaned Dataframe:\n", cleaned_df)

```

OUTPUT:

```

Original dataframe:
   numerical_col_1  numerical_col_2 categorical_col \
0                1.0             10.5             A
1                2.0             11.5             B
2                NaN             10.8             A
3                4.0              NaN             C
4                5.0             12.1             NaN

   sone_numeric_column_string  date_column  column_a  column_b
0                        100    2023-01-01         1         5
1                        200    2023-01-02         2         4
2                        abc  invalid date         3         3
3                        400    2023-01-04         4         2
4                        500    2023-01-05         5         1

-----Handling missing value-----
Missing values before cleaning:
numerical_col_1      1
numerical_col_2      1
categorical_col      1
sone_numeric_column_string  0
date_column          0
column_a             0
column_b             0
dtype: int64
Missing values after filling:
numerical_col_1      0
numerical_col_2      0
categorical_col      0
sone_numeric_column_string  0
date_column          0
column_a             0
column_b             0
dtype: int64

-----Type conversion-----
Converted 'sone_numeric_column_string' to numeric
Converted 'date_column' to datetime
-----Data transformation-----
Created new feature by multiplying 'column_a' and 'column_b'
-----Removing duplicates-----
Removed 0 duplicate rows

```

2.4) To detect outliers in the given data.

```
import pandas as pd

import numpy as np

InputFileName='Movie_collection_train.csv'

print('#####')

print("Input file")

sFileName='/content/Movie_collection_train.csv'

print('Loading :',sFileName)

Movie_DATA_ALL = pd.read_csv(sFileName, header=0, usecols=['Genre', '3D_available', 'Budget'],
encoding='latin-1')

Movie_DATA_ALL.rename(columns={'Genre':'Movie type'},inplace=True)

print(Movie_DATA_ALL)

MeanData=Movie_DATA_ALL.groupby(['Movie type','3D_available'])['Budget'].mean()

stdData=Movie_DATA_ALL.groupby(['Movie type','3D_available'])['Budget'].std()

print(MeanData);

print(stdData);

print('Outliers')

UpperBound = float(sum(MeanData) + sum(stdData))

print('Higher than ', UpperBound)

OutliersHigher = Movie_DATA_ALL[Movie_DATA_ALL.Budget > UpperBound]

print(OutliersHigher)

LowerBound = float(sum(MeanData) - sum(stdData))

print('Lower than ', LowerBound)

OutliersLower = Movie_DATA_ALL[Movie_DATA_ALL.Budget < LowerBound]

print(OutliersLower)

print('Not Outliers')

OutliersNot = Movie_DATA_ALL[(Movie_DATA_ALL.Budget > LowerBound) &
(Movie_DATA_ALL.Budget <= UpperBound)]

print(OutliersNot)
```

OUTPUT:

```
#####
Input file
Loading : /content/Movie_collection_train.csv
      Budget Movie type 3D_available
0      36524.125    Thriller      YES
1      35668.655      Drama      NO
2      39912.675    Comedy      NO
3      38873.890      Drama      YES
4      39701.585      Drama      NO
..      ...      ...      ...
395    35946.405      Action      NO
396    35579.775    Thriller      YES
397    31924.585    Comedy      NO
398    30291.415      Drama      NO
399    32507.860    Thriller      NO

[400 rows x 3 columns]
Movie type  3D_available
Action      NO      34097.978750
            YES      36832.983000
Comedy      NO      34607.650000
            YES      35025.761549
Drama      NO      33543.050588
            YES      36053.486489
Thriller    NO      34861.704453
            YES      35952.279253
Name: Budget, dtype: float64
Movie type  3D_available
Action      NO      2932.597166
            YES      4009.371358
Comedy      NO      3503.724780
            YES      3586.247444
Drama      NO      4649.035882
            YES      3655.962513
Thriller    NO      3998.570826
            YES      4754.690793
Name: Budget, dtype: float64
Outliers
Higher than  312065.0948452664
Empty DataFrame
Columns: [Budget, Movie type, 3D_available]
```



```
Columns: [Budget, Movie type, 3D_available]
Index: []
Lower than  249884.69332051632
      Budget Movie type 3D_available
0      36524.125    Thriller      YES
1      35668.655      Drama      NO
2      39912.675    Comedy      NO
3      38873.890      Drama      YES
4      39701.585      Drama      NO
..      ...      ...      ...
395    35946.405      Action      NO
396    35579.775    Thriller      YES
397    31924.585    Comedy      NO
398    30291.415      Drama      NO
399    32507.860    Thriller      NO

[400 rows x 3 columns]
Not Outliers
Empty DataFrame
Columns: [Budget, Movie type, 3D_available]
Index: []
```

Aim: 3.1) To perform regression analysis using single linear regression.

```
y = data_set.iloc[:, 1].values
import matplotlib.pyplot as mtp
import pandas as pd
from sklearn.model_selection import train_test_split
from sklearn.linear_model import LinearRegression
from sklearn.metrics import r2_score
# Load dataset
data_set = pd.read_csv("/content/Salary_Data - Salary_Data.csv")
# Independent variable (experience) and dependent variable (salary)
x = data_set.iloc[:, :- 1].values
y = data_set.iloc[:, 1].values
# Split into training and test sets
x_train, x_test, y_train, y_test = train_test_split( x, y, test_size=1/3,
random_state=0)
# Train the model
regressor = LinearRegression()
regressor.fit(x_train, y_train)
y_pred_train = regressor.predict(x_train)
mtp.scatter(x_train, y_train, color="green", label="Actual Salary (Training)")
mtp.plot(x_train, y_pred_train, color="red", label="Regression Line")
mtp.title("Salary vs Experience (Training Dataset)")
mtp.xlabel("Years of Experience")
mtp.ylabel("Salary (In Rupees)")
mtp.legend()
mtp.show()
# Predict test data
y_pred_test = regressor.predict(x_test)
# Find accuracy (R2 score)
accuracy = r2_score(y_test, y_pred_test)
print(f"Test Data Accuracy (R2 Score): {accuracy:.2f}")
# Plot actual vs predicted for test set
mtp.scatter(x_test, y_test, color="blue", label="Actual Salary")
mtp.scatter(x_test, y_pred_test, color="red", label="Predicted Salary")
mtp.plot( x_train, regressor.predict(x_train), color="green",
label="Regression Line")
mtp.title("Salary vs Experience (Test Dataset)")
mtp.xlabel("Years of Experience")
mtp.ylabel("Salary (In Rupees)")
mtp.legend()
mtp.show()
```

OUTPUT:



Aim:3.2) To perform regression analysis using multiple linear regression.

Aim:3.3) To perform logistic regression analysis

```
import pandas as pd
import numpy as np
import matplotlib.pyplot as plt
from sklearn.model_selection import train_test_split
from sklearn.linear_model import LogisticRegression
from sklearn.metrics import accuracy_score, confusion_matrix, classification_report
# Step 1: Create a dataset (Heart Disease like) as DataFrame
np.random.seed(42)
n_samples = 500
data = pd.DataFrame({
    "Age": np.random.randint( 29,77, n_samples),
    "Sex": np.random.randint( 0,2, n_samples), # 0 = female, 1 = male
    "Cholesterol": np.random.randint(150, 300, n_samples),
    "BloodPressure": np.random.randint( 90, 180, n_samples),
    "MaxHeartRate": np.random.randint( 90,200, n_samples)
})
# Target variable (rule-based: high Cholesterol, high BP, or low MaxHR + higher risk)
data["HeartDisease"] = ((data["Cholesterol"] > 240) |
    (data["BloodPressure"] > 140) |
    (data["MaxHeartRate"] < 120)).astype(int)
print("Sample of Heart Disease Dataset: \n")
print(data.head())
# Step 2: Split features & target
X = data.drop("HeartDisease",axis=1)
y = data["HeartDisease"]
X_train, X_test, y_train, y_test = train_test_split( X, y, test_size=0.3, random_state=42)
# Step 3: Train Logistic Regression
model = LogisticRegression(max_iter=500)
model.fit(X_train, y_train)
# Step 4: Predictions & Evaluation
y_pred = model.predict(X_test)
print("\n Model Evaluation:")
print("Accuracy:", accuracy_score(y_test, y_pred))
print("\nConfusion Matrix:\n", confusion_matrix(y_test, y_pred))
print("\nClassification Report:\n", classification_report(y_test,y_pred))
OUTPUT:
```


Sample of Heart Disease Dataset:

	Age	Sex	Cholesterol	BloodPressure	MaxHeartRate	HeartDisease
0	67	1	298	115	104	1
1	57	1	219	115	186	0
2	43	1	150	136	188	0
3	71	1	282	121	115	1
4	36	0	161	99	97	1

Model Evaluation:
Accuracy: 0.8533333333333334

Confusion Matrix:
[[37 11]
[11 91]]

Classification Report:

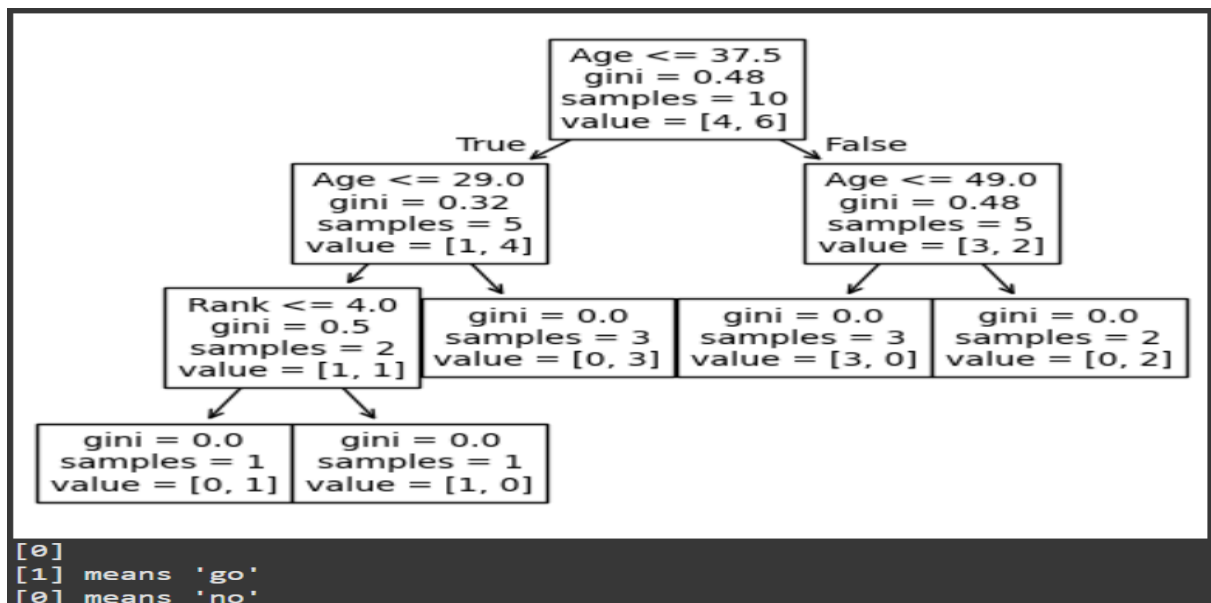
	precision	recall	f1-score	support
0	0.77	0.77	0.77	48
1	0.89	0.89	0.89	102
accuracy			0.85	150
macro avg	0.83	0.83	0.83	150
weighted avg	0.85	0.85	0.85	150

4.1) To implement classification using decision tree induction

```
import pandas as pd
import sys
from sklearn import tree
from sklearn.tree import DecisionTreeClassifier
import matplotlib.pyplot as plt
import pandas as pd
data = {'Age': [30, 45, 25, 35, 40, 50, 28, 32, 48, 55],
        'Experience': [5, 15, 2, 10, 12, 20, 4, 7, 18, 25],
        'Rank': [7, 9, 3, 8, 7, 10, 5, 6, 9, 10],
        'Nationality': ['UK', 'USA', 'N', 'UK', 'USA', 'UK', 'N', 'USA', 'UK', 'N'],
        'Go': ['YES', 'NO', 'YES', 'YES', 'NO', 'YES', 'NO', 'YES', 'NO', 'YES']}

df = pd.DataFrame(data)
df.to_csv('/content/dataset.csv', index=False)
d={'UK':0,'USA':1,'N':2}
df['Nationality']=df['Nationality'].map(d)
d={'YES':1,'NO':0}
df['Go']=df['Go'].map(d)
fea=['Age','Experience','Rank','Nationality']
x=df[fea]
y=df['Go']
dtree=DecisionTreeClassifier()
dtree=dtree.fit(x,y)
tree.plot_tree(dtree,feature_names=fea)
plt.savefig("result.png")
plt.show()
print(dtree.predict([[40,10,7,1]]))
print("[1] means 'go'")
print("[0] means 'no'")
```

OUTPUT:



4.2) To implement classification using Naïve Bayes algorithm

```

import pandas as pd

from sklearn.model_selection import train_test_split
from sklearn.preprocessing import LabelEncoder
from sklearn.naive_bayes import GaussianNB
from sklearn.metrics import accuracy_score, confusion_matrix, classification_report

# Load the dataset

data = pd.read_csv('/content/loan.csv')

# View the first few rows

print(data.head())

# Drop rows with missing values (you can also choose to impute)

data.dropna(inplace=True)

# Encode categorical variables directly in original DataFrame to avoid SettingWithCopyWarning

le = LabelEncoder()

for col in ['Gender', 'Married', 'Education', 'Self_Employed']:

    data.loc[:, col] = le.fit_transform(data[col])

# Encode target variable

data['Loan_Status'] = le.fit_transform(data['Loan_Status'])

# Select features and target after

X = data[['Gender', 'Married', 'Education', 'Self_Employed', 'ApplicantIncome', 'LoanAmount']]
  
```

```

y = data['Loan_Status' ]

# Split into train and test sets

X_train, X_test, y_train, y_test = train_test_split( X, y, test_size=0.3, random_state=42)

# Initialize and train Naive Bayes model

model = GaussianNB()

model.fit(X_train, y_train)

#Predict on test set

y_pred = model.predict(X_test)

#Evaluate the model

print("Accuracy:", accuracy_score(y_test, y_pred))

print("Confusion Matrix:\n", confusion_matrix(y_test, y_pred))

print("Classification Report:\n", classification_report(y_test, y_pred))

# Create a sample DataFrame with the correct feature names and order

sample = pd.DataFrame ({

'Gender': [1], # Male encoded as 1

'Married': [1], # Yes encoded as 1

'Education': [1], # Graduate encoded as 1

'Self_Employed': [0], # No encoded as 0

'ApplicantIncome': [5000],

'LoanAmount': [128]

})

# Predict the class for the unknown sample or evidence

predicted_class = model.predict(sample)

# Map prediction back to label

loan_status_map = {0: 'N', 1: 'Y'}

print(f"Predicted Loan Status: {loan_status_map[predicted_class[0]]}")

```

OUTPUT:

```

Loan_ID Gender Married Dependents Education Self_Employed \
0 LP001002 Male No 0 Graduate No
1 LP001003 Male Yes 1 Graduate No
2 LP001005 Male Yes 0 Graduate Yes
3 LP001006 Male Yes 0 Not Graduate No
4 LP001008 Male No 0 Graduate No

ApplicantIncome CoapplicantIncome LoanAmount Loan_Amount_Term \
0 5849 0.0 NaN 360.0
1 4583 1508.0 128.0 360.0
2 3000 0.0 66.0 360.0
3 2583 2358.0 120.0 360.0
4 6000 0.0 141.0 360.0

Credit_History Property_Area Loan_Status
0 1.0 Urban Y
1 1.0 Rural N
2 1.0 Urban Y
3 1.0 Urban Y
4 1.0 Urban Y
Accuracy: 0.6875
Confusion Matrix:
[[ 3 41]
 [ 4 96]]
Classification Report:
precision recall f1-score support
0 0.43 0.07 0.12 44
1 0.70 0.96 0.81 100

accuracy 0.69 144
macro avg 0.56 0.51 0.46 144
weighted avg 0.62 0.69 0.60 144

Predicted Loan Status: Y

```

Aim:4.3) To implement classification using decision tree induction with various attribute selection methods(Information Gain, Gini index and Gain ratio)

Aim: 5.1 To implement clustering using K-Means Algorithm

```
import pandas as pd
import matplotlib.pyplot as plt
from sklearn.cluster import KMeans
from sklearn.preprocessing import StandardScaler

# Step 1: Synthetic cluster-friendly dataset
data = {
    'Age': [22, 23, 25, 24, 26,      # Group 1: Young, low salary, short browsing
            35, 36, 34, 33, 37,      # Group 2: Mid-age, mid salary, medium browsing
            48, 50, 52, 49, 51,      # Group 3: Older, high salary, long browsing
            23, 36, 50, 35, 48],     # Mix for variation
    'Salary': [25000, 27000, 26000, 28000, 24000,
               60000, 62000, 58000, 61000, 59000,
               100000, 98000, 105000, 97000, 102000,
               25500, 60500, 101000, 61500, 99000],
    'Browsing_Time': [1.5, 1.8, 2.0, 1.6, 1.9,
                      5.0, 5.2, 4.8, 5.5, 5.1,
                      9.0, 8.5, 9.2, 8.8, 9.5,
                      2.0, 5.3, 9.0, 5.0, 8.7]
}

df = pd.DataFrame(data)
```

```

# Step 2: Scale features
scaler = StandardScaler()
scaled = scaler.fit_transform(df)


# Step 3: KMeans clustering
kmeans = KMeans(n_clusters=3, random_state=42)
df['Cluster'] = kmeans.fit_predict(scaled)

# Step 4: Plot clusters (Age vs Salary)
plt.figure(figsize=(8, 6))
for cluster in df['Cluster'].unique():
    cluster_data = df[df['Cluster'] == cluster]
    plt.scatter(cluster_data['Age'], cluster_data['Browsing_Time'], label=f'Cluster {cluster}', s=100)

plt.title('Clustered Data (Age vs Salary)')
plt.xlabel('Age')
plt.ylabel('Browsing_Time')
plt.legend()
plt.grid(True)
plt.show()

```

OUTPUT:



Aim: 5.2) To perform hierarchical clustering

```
import pandas as pd
import matplotlib.pyplot as plt
import seaborn as sns
from scipy.cluster.hierarchy import linkage, dendrogram
from sklearn.preprocessing import StandardScaler
from sklearn.cluster import AgglomerativeClustering

#Step1:
data={
    'Age':[22,25,47,52,48,55,60,32,44,25,
           40,28,38,29,30,41,26,34,45,50],
    'Salary':[25000,27000,90000,110000,95000,120000,99000,105000,115000,48000,
              80000,30000,75000,32000,35000,82000,28000,60000,87000,100000],
    'Browsing_Time':[1.5,2.0,8.5,9.0,7.5,10.0,7.0,8.0,9.5,3.5,
                     6.5,2.5,6.0,3.0,3.2,7.0,2.2,4.5,6.8,8.5]
}
df=pd.DataFrame(data)

#Step2
scaler=StandardScaler()
X_scaled=scaler.fit_transform(df)

#step3
plt.figure(figsize=(10,6))
```

```


linked=linkage(X_scaled,method='ward')
dendrogram(linked,
            orientation='top',
            distance_sort='ascending',
            show_leaf_counts=True)
plt.title("Dendrogram")
plt.xlabel("Sample")
plt.ylabel("Distance")
plt.show()

#step4
cluster=AgglomerativeClustering(n_clusters=4,linkage='ward')
df['Cluster']=cluster.fit_predict(X_scaled)

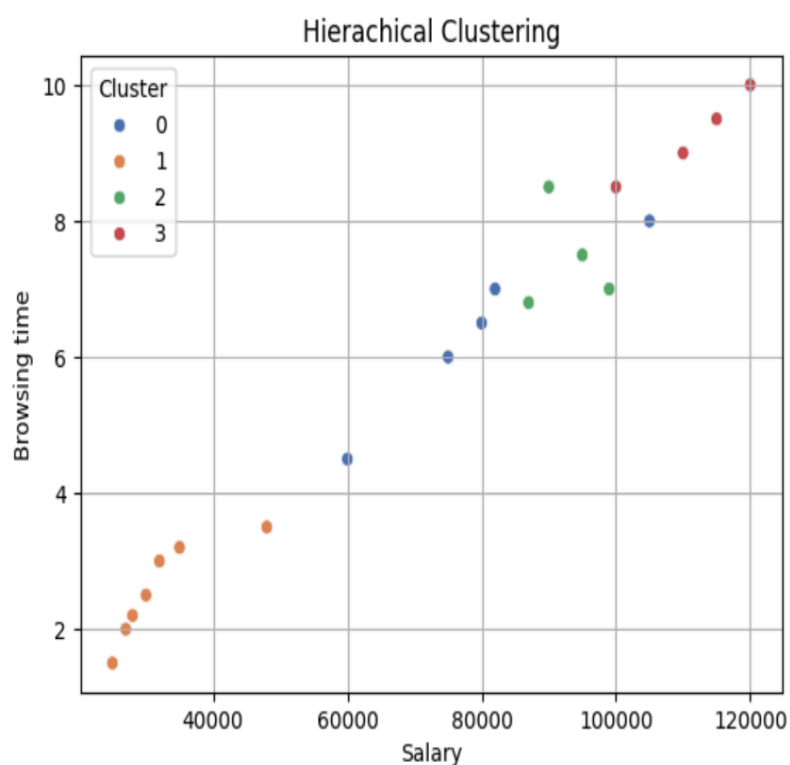
#step5
print("Clustered Data")
print(df)
sns.scatterplot(data=df,x='Salary',y='Browsing_Time',hue='Cluster',palette='deep')
plt.title("Hierarchical Clustering")
plt.xlabel("Salary")
plt.ylabel("Browsing time")
plt.grid(True)
plt.show()

```

OUTPUT:



Clustered	Date	Age	Salary	Browsing_Time	Cluster
0		22	25000	1.5	1
1		25	27000	2.0	1
2		47	90000	8.5	2
3		52	110000	9.0	3
4		48	95000	7.5	2
5		55	120000	10.0	3
6		60	99000	7.0	2
7		32	105000	8.0	0
8		44	115000	9.5	3
9		25	48000	3.5	1
10		40	80000	6.5	0
11		28	30000	2.5	1
12		38	75000	6.0	0
13		29	32000	3.0	1
14		30	35000	3.2	1
15		41	82000	7.0	0
16		26	28000	2.2	1
17		34	60000	4.5	0
18		45	87000	6.8	2
19		50	100000	8.5	3



Practical 6: To implement PCA (Principal Component Analysis).

```
import pandas as pd
```

```
import numpy as np
```

```
import matplotlib.pyplot as plt
```

```
from sklearn.decomposition import PCA
```

```
from sklearn.preprocessing import StandardScaler
```

```
from sklearn.datasets import load_breast_cancer
```

```
data=load_breast_cancer()
```

```
data.keys()
```

```

print(data['target_names']) #check the output class
print(data['feature_names']) #check the input features
df1=pd.DataFrame(data['data'],columns=data['feature_names'])
scaling=StandardScaler()
scaling.fit(df1)
scaled_data=scaling.transform(df1)
principal=PCA(n_components=3) #set n_components=3
principal.fit(scaled_data)
x=principal.transform(scaled_data)
print(x.shape)
plt.figure(figsize=(10,10))
plt.scatter(x[:,0],x[:,1],c=data['target'],cmap='plasma')
plt.xlabel('pc1')
plt.ylabel('pc2')
plt.show()

from mpl_toolkits.mplot3d import Axes3D
fig=plt.figure(figsize=(10,10))
axis=fig.add_subplot(111,projection='3d')
axis.scatter(x[:,0],x[:,1],x[:,2],c=data['target'],cmap='plasma')

axis.set_xlabel('pc1',fontsize=10)
axis.set_ylabel('pc2',fontsize=10)
axis.set_zlabel('pc3',fontsize=10)
plt.show()



```

OUTPUT:

```

['malignant' 'benign']
['mean radius' 'mean texture' 'mean perimeter' 'mean area'
 'mean smoothness' 'mean compactness' 'mean concavity'
 'mean concave points' 'mean symmetry' 'mean fractal dimension'
 'radius error' 'texture error' 'perimeter error' 'area error'
 'smoothness error' 'compactness error' 'concavity error'
 'concave points error' 'symmetry error' 'fractal dimension error'
 'worst radius' 'worst texture' 'worst perimeter' 'worst area'
 'worst smoothness' 'worst compactness' 'worst concavity'
 'worst concave points' 'worst symmetry' 'worst fractal dimension']
(569, 3)

```



Practical 7 :To explore the given data and identify the patterns in it.

Aim: 8.1) To evaluate binary classification model using confusion matrix along with precision and recall.


```
from sklearn.datasets import load_breast_cancer
from sklearn.model_selection import train_test_split
from sklearn.tree import DecisionTreeClassifier
from sklearn.metrics import confusion_matrix
```

```

import seaborn as sns
import matplotlib.pyplot as plt
from sklearn.metrics import accuracy_score, precision_score, recall_score, f1_score
#load dataset
X,y=load_breast_cancer(return_X_y=True)
X_train,X_test,y_train,y_test=train_test_split(X,y,test_size=0.25)
#Train model
tree=DecisionTreeClassifier(random_state=23)
tree.fit(X_train,y_train)
y_pred=tree.predict(X_test)
#compute the matrix
cm=confusion_matrix(y_test,y_pred)
#plot the matrix
sns.heatmap(cm,
            annot=True,
            fmt='g',
            xticklabels=['malignant','benign'],
            yticklabels=['malignant','benign'])
plt.ylabel("prediction")
plt.xlabel("actual")
plt.title("confusion matrix")
plt.show()
#finding presion and recall
accuracy=accuracy_score(y_test,y_pred)
print("accuracy:",accuracy)
precision=precision_score(y_test,y_pred)
print("presion:",precision)
recall=recall_score(y_test,y_pred)
print("recall:",recall)
f=f1_score(y_test,y_test)
print("f1:",f)

```

OUTPUT:



Aim:8.2) To evaluate multi-class classification model using confusion matrix along with precision and recall.

Aim:9(Use an appropriate dataset and create a supervised learning model, Analyse the model with ROC-AUC.

#Use an appropriate dataset and create a supervised learning model, Analyse the model with ROC-AUC.

```
import pandas as pd
```

```
import seaborn as sns
```

```

import matplotlib.pyplot as plt

from sklearn.tree import DecisionTreeClassifier

from sklearn.datasets import load_breast_cancer

from sklearn.model_selection import train_test_split

from sklearn.metrics import (
    confusion_matrix, accuracy_score, precision_score,
    recall_score, f1_score, roc_curve, auc)

X, y = load_breast_cancer(return_X_y=True)

X_train, X_test, y_train, y_test = train_test_split(
    X, y, test_size=0.25, random_state=23, stratify=y
)

tree1 = DecisionTreeClassifier(random_state=23)

tree1.fit(X_train, y_train)

y_pred1 = tree1.predict(X_test)

y_proba1 = tree1.predict_proba(X_test)[:, 1]

print("=== Breast Cancer Dataset ===")

print("Accuracy :", accuracy_score(y_test, y_pred1))

print("Precision:", precision_score(y_test, y_pred1))

print("Recall  :", recall_score(y_test, y_pred1))

print("F1-score :", f1_score(y_test, y_pred1))


cm1 = confusion_matrix(y_test, y_pred1)

sns.heatmap(cm1, annot=True, fmt="g",
            xticklabels=["malignant", "benign"],
            yticklabels=["malignant", "benign"])

plt.title("Confusion Matrix - Breast Cancer Dataset")

plt.show()

fpr1, tpr1, _ = roc_curve(y_test, y_proba1)

roc_auc1 = auc(fpr1, tpr1)


data = pd.DataFrame({

```

```

    "education":
["bach","mast","diploma","mast","diploma","bach","mast","mast","diploma","mast","bach","mast",
"mast","mast","bach"],

    "job": [1,0,0,0,0,0,0,0,0,1,0,0,0,0]

})

```

```

# Encode categorical features

```

```

X2 = pd.get_dummies(data.drop("job", axis=1))
y2 = data["job"]

```

```

X2_train, X2_test, y2_train, y2_test = train_test_split(
    X2, y2, test_size=0.25, random_state=23, stratify=y2
)

```

```

tree2 = DecisionTreeClassifier(random_state=23)
tree2.fit(X2_train, y2_train)

```

```

y2_pred = tree2.predict(X2_test)
y2_proba = tree2.predict_proba(X2_test)[:, 1]

```

```

print("\n=== Unbalanced CSV Dataset ===")
print("Accuracy :", accuracy_score(y2_test, y2_pred))
print("Precision:", precision_score(y2_test, y2_pred))
print("Recall  :", recall_score(y2_test, y2_pred))
print("F1-score :", f1_score(y2_test, y2_pred))

```

```

cm2 = confusion_matrix(y2_test, y2_pred)
sns.heatmap(cm2, annot=True, fmt="g")
plt.title("Confusion Matrix - Unbalanced CSV Dataset")
plt.show()

```

```

fpr2, tpr2, _ = roc_curve(y2_test, y2_proba)

```



```
roc_auc2 = auc(fpr2, tpr2)
```

```
plt.figure(figsize=(7, 5))
```

```
plt.plot(fpr1, tpr1, label=f"Breast Cancer (AUC={roc_auc1:.2f})")
```

```
plt.plot(fpr2, tpr2, label=f"Unbalanced CSV (AUC={roc_auc2:.2f})")
```

```
plt.plot([0, 1], [0, 1], "r--", label="Random Guess")
```

```
plt.xlabel("False Positive Rate")
```

```
plt.ylabel("True Positive Rate")
```

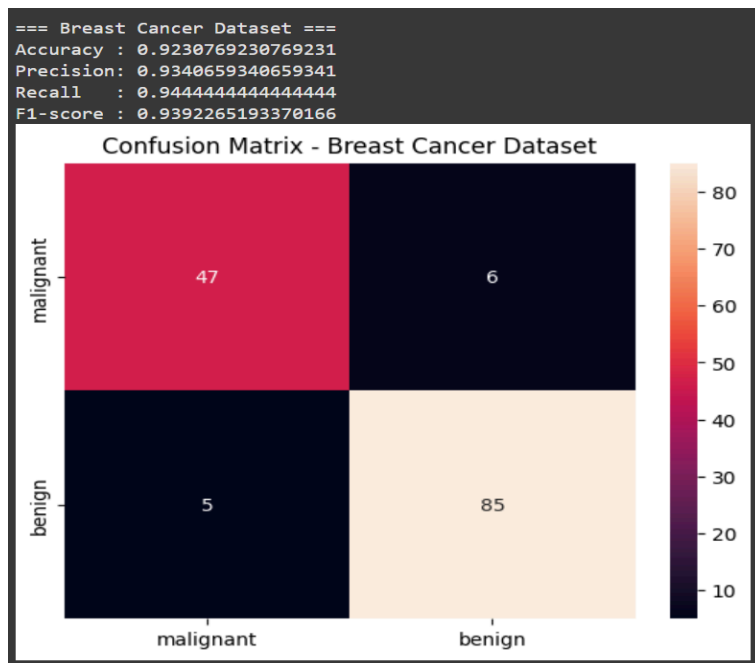
```
plt.title("ROC Curves Comparison")
```

```
plt.legend()
```

```
plt.grid(True)
```

```
plt.show()
```

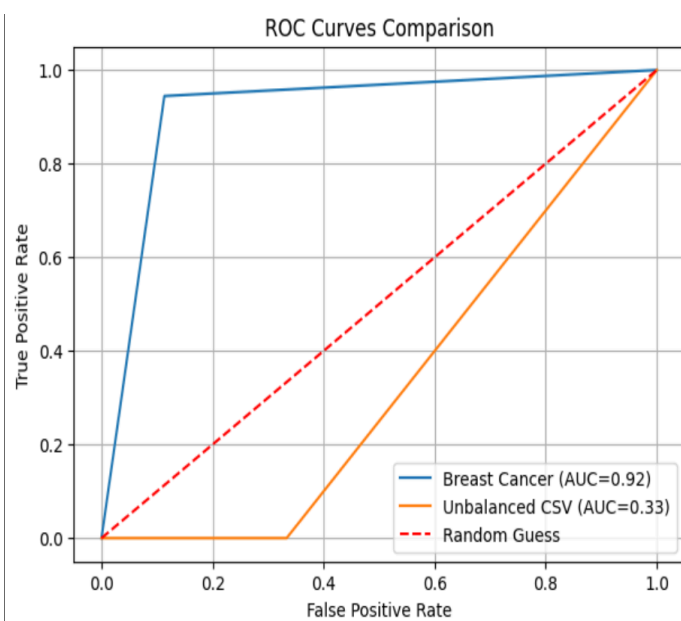

OUTPUT:



```

=== Unbalanced CSV Dataset ===
Accuracy : 0.75
Precision: 0.0
Recall   : 0.0
F1-score : 0.0
/usr/local/lib/python3.12/dist-packages/sklearn/metrics/_classification.py:1465: UserWarning:
  _warn_prf(average, modifier, f"{metric.capitalize()} is", len(res

```



Practical 10.:Consider a case study problem and implement an appropriate model and evaluate it.

Aim: 11.1)Bagging and boosting model.

```
#bagging and boosting

from sklearn.datasets import load_breast_cancer
from sklearn.model_selection import train_test_split
from sklearn.ensemble import RandomForestClassifier,GradientBoostingClassifier
from sklearn.metrics import classification_report
import pandas as pd

data = pd.read_csv("/content/breast-cancer.csv")
x=data.drop("diagnosis", axis=1)
y=data.diagnosis
df=pd.DataFrame(y)
print(df.head())

#splitting data
x_train,x_test,y_train,y_test = train_test_split(x,y,test_size=0.2,random_state=42)

#initialize model
rf=RandomForestClassifier(n_estimators=100,random_state=42)
gb=GradientBoostingClassifier(n_estimators=100,learning_rate=0.3,random_state=42)

#train model
rf.fit(x_train,y_train)
gb.fit(x_train,y_train)

#predict model
y_pred_rf=rf.predict(x_test)
y_pred_gb=gb.predict(x_test)

#Evaluate and print result
print("\nRandom forest (bagging) classification report: ")
print(classification_report(y_test,y_pred_rf))
print("\nGradient boosting (boosting) classification report: ")
print(classification_report(y_test,y_pred_gb))
```

OUTPUT:

```
diagnosis
0      M
1      M
2      M
3      M
4      M

Random forest (bagging) classification report:
              precision    recall  f1-score   support

     B       0.96       0.99       0.97       71
     M       0.98       0.93       0.95       43

 accuracy          0.96          114
 macro avg       0.97       0.96       0.96       114
weighted avg       0.97       0.96       0.96       114

Gradient boosting (boosting) classification report:
              precision    recall  f1-score   support

     B       0.96       0.96       0.96       71
     M       0.93       0.93       0.93       43

 accuracy          0.95          114
 macro avg       0.94       0.94       0.94       114
weighted avg       0.95       0.95       0.95       114
```

Aim:11.2 Cross validation methods

```
from sklearn.datasets import load_breast_cancer
from sklearn.ensemble import RandomForestClassifier, GradientBoostingClassifier
from sklearn.model_selection import cross_val_score, StratifiedKFold
import numpy as np
data = load_breast_cancer()
X = data.data
y = data.target
kf = StratifiedKFold(n_splits=10)
rf = RandomForestClassifier(n_estimators=100, n_jobs=42)
gb = GradientBoostingClassifier(n_estimators=100, learning_rate=0.1, random_state=42)
rf_scores = cross_val_score(rf, X, y, cv=kf, scoring='f1')
gb_scores = cross_val_score(gb, X, y, cv=kf, scoring='f1')
print(f"Random forest (Bagging) 10-fold cv f1-score:"
      f"Mean={rf_scores.mean():.4f}; Std={rf_scores.std():.4f}")
print(f"Gradient boosting (Boosting) 10-fold cv f1-score:"
      f"Mean={gb_scores.mean():.4f}; Std={gb_scores.std():.4f}")
```

OUTPUT:

```
Random forest (Bagging) 10-fold cv f1-score:Mean=0.9724;,Std=0.0192
Gradient boosting (Boosting) 10-fold cv f1-score:Mean=0.9694;,Std=0.0235
```